

STATS 790

Final Project

Dean Hansen - 400027416

28 April, 2023

Contents

Introduction	2
Random Forest Algorithm	3
Califonia Housing Dataset	4
Adults Dataset	5
randomForest Approach	6
Scikit Learn Approach	7
Fitted Model Comparisons	8
Conclusions	9
References	10

Introduction

In this project, we will compare the R randomForest and Python sci-kit learn implementations of the random forest algorithm. We will fit a regression and classification model using both libraries and compare their performance and various metrics of speed and prediction accuracy.

Random Forest Algorithm

The random forest algorithm was created by Leo Breiman, and can be thought of as an improvement to the bagging algorithm (Breiman 2001). One key difference between random forests and bagging is when each tree is grown, only a random subset of features are used at each terminal node to split the data. Whereas in bagging, the full set of features are used at each binary split. The purpose of this change is to decrease the correlation among the collection of trees. Below we present the random forest algorithm found in the Elements of Statistical Learning (Second Edition) on page 588 (Hastie, Tibshirani, and Friedman 2009).

Algorithm 15.1 *Random Forest for Regression or Classification.*

1. For $b = 1$ to B :
 - (a) Draw a bootstrap sample \mathbf{Z}^* of size N from the training data.
 - (b) Grow a random-forest tree T_b to the bootstrapped data, by recursively repeating the following steps for each terminal node of the tree, until the minimum node size n_{min} is reached.
 - i. Select m variables at random from the p variables.
 - ii. Pick the best variable/split-point among the m .
 - iii. Split the node into two daughter nodes.
2. Output the ensemble of trees $\{T_b\}_1^B$.

To make a prediction at a new point x :

Regression: $\hat{f}_{\text{rf}}^B(x) = \frac{1}{B} \sum_{b=1}^B T_b(x)$.

Classification: Let $\hat{C}_b(x)$ be the class prediction of the b th random-forest tree. Then $\hat{C}_{\text{rf}}^B(x) = \text{majority vote } \{\hat{C}_b(x)\}_1^B$.

Figure 1: Random Forest algorithm

Note: predictions are made by majority vote (classification) or by averaging the predicted value from each tree (regression).

California Housing Dataset

The **California housing** dataset, found in section 10.14.1 of ESL, comes from the paper on “Sparse spatial autoregressions” (Pace and Barry 1997). There are 20,640 observations, 8 predictor variables and one response variable, median house value. A detailed description of each variable can be found below.

Variable	Description
median_income	median income of California block
housing_median_age	median age of the house
total_rooms	aggregate total number of rooms
total_bedrooms	aggregate total number of bedrooms
population	population of geographic block
households	number of houses
latitude	latitude of geographic block
longitude	longitude of geographic block
median_house_value (response*)	median house value

We downloaded the California housing dataset from the Statlib website found [here](#). The Random Forest algorithm does not require feature scaling, so no data pre-processing is required.

Adults Dataset

The **Adults** dataset comes from the UCI Machine Learning Repository (Kohavi and Becker 1996). The dataset contains 45,170 observations and 15 features with information such as age, sex, hours of work per week which are used to classify whether an individual earns more than 50,000 USD per year. The original dataset was sourced by Barry Becker from the US 1994 Census database. Further, a detailed description of each variable can be found below (description sourced from (Eren 2018)).

Variable	Description
age	Person's age
workclass	Occupation sector (ex. private, federal government)
finalweight	Estimate of population total where the individual is sourced from (i.e. similar observations should have similar weights based on Census data)
education	Highest level of education attained (ex. bachelors)
educationnum	Years of education
maritalstatus	Marital status (ex. divorced, widowed)
occupation	Occupation title (ex. tech-support, farming-fishing)
relationship	Relationship status (ex. wife, unmarried)
race	Person's race (ex. white, asian-pac-islander)
sex	Male or Female
capitalgain	Capital gains from the previous year
capitalloss	Capital losses from the previous year
hoursperweek	Number of hours worked per week
nativecountry	Country person was born (ex. Ireland, Cuba)
isGT50K (response*)	1 if they earn more than 50,000 USD per year and 0 otherwise

randomForest Approach

First we create a random forest regression model using the California Housing dataset. The `randomForest` package can be used as the engine, so we will do it this way to make it easier to setup a tuning grid.

Due to computing constraints, a small grid of values have been selected to make the algorithm comparisons.

We adjust the above code and set the mode to classification, this time using the Adults dataset.

Scikit Learn Approach

There are slight variations between the R and Python implementations, so we have tried to make them as similar as possible so a valid comparison can be made. We use the Scikit Learn library implementation of random forest to fit the regression and classification model (Pedregosa et al. 2011).

The `randomforestclassifier` did not accept the original dataset with strings, so it was converted using pandas to a set of dummy variables first using help from [here](#).

Fitted Model Comparisons

Using the same default settings, randomForest had an accuracy of 86.62% compared to 85.23% for sci-kit learn for classification. For regression, sci-kit learn had a RMSE of 49397.93 compared to 50829.69 for randomForest.

Further, sci-kit learn was significantly faster than the R implementation.

Conclusions

From the comparisons, we see that the Python implementation is generally faster compared to the R approach. There are some potential pitfalls in this analysis that cause the outputted times to vary, this could be due to lack of parallelization, variation in internals of each algorithm and hyperparameter variation. We have tried to reduce the last piece by tuning both in R and Python first, then creating a final model that uses the same hyperparameter values.

The out-of-bag error plots demonstrate the two approaches to be quite similar. Further, the predictions on the validation sets for both classification and regression appear to be similar up to a high order of floating point approximation.

In conclusion, the choice between Python or R comes down to user preference.

References

- Breiman, Leo. 2001. *Machine Learning* 45 (1): 5–32. <https://doi.org/10.1023/a:1010933404324>.
- Eren, Kazim Anil. 2018. “Adult Census Income Classification.” Kaggle. <https://www.kaggle.com/code/kazimanil/adult-census-income-classification>.
- Hastie, Trevor, Robert Tibshirani, and Jerome Friedman. 2009. *The Elements of Statistical Learning: Data Mining, Inference and Prediction*. 2nd ed. Springer. <http://www-stat.stanford.edu/~tibs/ElemStatLearn/>.
- Kohavi, Ronny, and Barry Becker. 1996. *UCI Machine Learning Repository*. University of California, Irvine, School of Information; Computer Sciences. <http://archive.ics.uci.edu/ml/datasets/Adult>.
- Pace, R. Kelley, and Ronald Barry. 1997. “Sparse Spatial Autoregressions.” *Statistics and Probability Letters* 33 (3): 291–97. [https://doi.org/10.1016/s0167-7152\(96\)00140-x](https://doi.org/10.1016/s0167-7152(96)00140-x).
- Pedregosa, F., G. Varoquaux, A. Gramfort, V. Michel, B. Thirion, O. Grisel, M. Blondel, et al. 2011. “Scikit-Learn: Machine Learning in Python.” *Journal of Machine Learning Research* 12: 2825–30.